

dm-1-asst-3

November 25, 2023

```
[1]: import pandas as pd
import seaborn as sns
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules
```

1.Association Rule Generation from Transaction Data

```
[2]: data = pd.read_csv('/kaggle/input/dataset/Grocery_Items_14.csv')
```

```
[3]: Transaction = []
for i in range(data.shape[0]):
    Transaction.append([str(data.values[i,j]) for j in range(data.shape[1]) if
↳str(data.values[i,j]) != 'nan'])
```

```
[4]: te = TransactionEncoder()
te_ary = te.fit(Transaction).transform(Transaction)
df = pd.DataFrame(te_ary, columns=te.columns_)
```

(C)

```
[5]: frequent_itemsets = apriori(df, min_support=0.01, use_colnames=True)
rules = association_rules(frequent_itemsets, metric="confidence",
↳min_threshold=0.1)
```

```
[6]: rules
```

```
[6]:
```

	antecedents	consequents	antecedent support	\
0	(soda)	(other vegetables)	0.097250	
1	(other vegetables)	(whole milk)	0.121250	
2	(rolls/buns)	(whole milk)	0.109625	
3	(soda)	(whole milk)	0.097250	
4	(yogurt)	(whole milk)	0.085750	

	consequent support	support	confidence	lift	leverage	conviction	\
0	0.121250	0.010250	0.105398	0.869266	-0.001542	0.982281	
1	0.157875	0.014625	0.120619	0.764013	-0.004517	0.957633	
2	0.157875	0.014500	0.132269	0.837809	-0.002807	0.970491	
3	0.157875	0.011125	0.114396	0.724598	-0.004228	0.950905	

```
4          0.157875  0.011625  0.135569  0.858708 -0.001913  0.974195
```

```
zhangs_metric
0      -0.142807
1      -0.260080
2      -0.178594
3      -0.296280
4      -0.152523
```

(D)

```
[7]: msv = [0.001,0.005,0.01,0.05]
```

```
mct = [0.05,0.075,0.1]
```

```
[8]: h_data = []
for i in msv:
    for j in mct:
        frequent_itemsets = apriori(df, min_support=i, use_colnames=True)
        rules = association_rules(frequent_itemsets, metric="confidence",
        min_threshold=j)
        res = len(rules)
        h_data.append({'msv': i, 'mct': j, 'rules_count': res})
```


```
[9]: heatdata = pd.DataFrame(h_data)
heatdata = heatdata.pivot(index='mct', columns='msv', values='rules_count')
```

```
[10]: heatdata
```

```
[10]: msv      0.001  0.005  0.010  0.050
mct
0.050      527     60     12      0
0.075      292     43     10      0
0.100      144     22      5      0
```

```
[11]: sns.heatmap(heatdata.sort_index(ascending=False))
```

```
[11]: <Axes: xlabel='msv', ylabel='mct'>
```



(E)

```
[12]: frequent_itemsets_2 = apriori(df, min_support=0.005, use_colnames=True)
      rules_2 = association_rules(frequent_itemsets_2, metric="confidence",
      ↪min_threshold=0.0)
```

```
[13]: rules_2.nlargest(1, 'confidence')
```

```
[13]:
```

	antecedents	consequents	antecedent support	consequent support	\
18	(frankfurter)	(other vegetables)	0.038125	0.12125	

	support	confidence	lift	leverage	conviction	zhangs_metric
18	0.006125	0.160656	1.324996	0.001502	1.046948	0.255003

2. Image Classification using CNN

```
[14]: import os
      import keras
      import matplotlib.pyplot as plt
      from sklearn.model_selection import train_test_split
      from tensorflow.keras.layers import Conv2D, Dense, MaxPool2D, Flatten, Dropout,
      ↪BatchNormalization, DepthwiseConv2D, InputLayer, GlobalAveragePooling2D
```

```
from tensorflow.keras.regularizers import l2
from tensorflow.keras.models import Sequential
```

```
[15]: directory = "/kaggle/input/reena-dataset"
```

```
[16]: def images_labels(data_directory):
    Images_data = []
    labels_data = []
    classes = os.listdir(data_directory)
    for i in classes:
        for j in os.listdir(os.path.join(data_directory,i)):
            Images_data.append(os.path.join(data_directory,i,j))
            labels_data.append(i)
    return([Images_data,labels_data])
```

```
[17]: def data_split(data,labels,train_size):
    train_imgs,test_imgs,train_label,test_label = train_test_split(data,
                                                                    labels,
                                                                    random_state=10)

    return([train_imgs,test_imgs,train_label,test_label])
```

```
[18]: Images_data,labels_data = images_labels(directory)
train_imgs,test_imgs,train_label,test_label = data_split(Images_data,labels_data,0.8)
```

```
[19]: datagen = keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
```

```
[20]: train_data = datagen.flow_from_dataframe(
    pd.DataFrame({'file_paths': train_imgs, 'labels': train_label}),
    x_col='file_paths',
    y_col='labels',
    target_size=(256, 256),
    batch_size=32,
    class_mode='categorical',
)
```

Found 527 validated image filenames belonging to 4 classes.

```
[21]: test_data = datagen.flow_from_dataframe(
    pd.DataFrame({'file_paths': test_imgs, 'labels': test_label}),
    x_col='file_paths',
    y_col='labels',
    target_size=(256, 256),
    batch_size=32,
    class_mode='categorical',
)
```

```
)
```

Found 132 validated image filenames belonging to 4 classes.

```
[22]: cnn_1 = Sequential([
        Conv2D(8, (3, 3), activation='relu', input_shape=(256, 256, 3)),
        MaxPool2D(2, 2),
        Flatten(),
        Dense(16, activation='relu'),
        Dense(4, activation='softmax')
    ])
cnn_1.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
cnn_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 8)	224
max_pooling2d (MaxPooling2D)	(None, 127, 127, 8)	0
flatten (Flatten)	(None, 129032)	0
dense (Dense)	(None, 16)	2064528
dense_1 (Dense)	(None, 4)	68

=====
Total params: 2064820 (7.88 MB)
Trainable params: 2064820 (7.88 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```
[23]: train_1 = cnn_1.fit(train_data,
                        validation_data = test_data,
                        epochs = 20)
```

Epoch 1/20

17/17 [=====] - 9s 472ms/step - loss: 1.9831 - accuracy: 0.2448 - val_loss: 1.4567 - val_accuracy: 0.2424

Epoch 2/20

17/17 [=====] - 8s 440ms/step - loss: 1.3392 - accuracy: 0.3795 - val_loss: 1.3668 - val_accuracy: 0.2652


Epoch 3/20

17/17 [=====] - 7s 412ms/step - loss: 1.1690 -
accuracy: 0.5256 - val_loss: 1.2282 - val_accuracy: 0.3939
Epoch 4/20
17/17 [=====] - 7s 417ms/step - loss: 0.9513 -
accuracy: 0.6224 - val_loss: 1.1976 - val_accuracy: 0.5152
Epoch 5/20
17/17 [=====] - 7s 403ms/step - loss: 0.7637 -
accuracy: 0.6907 - val_loss: 1.3209 - val_accuracy: 0.4773
Epoch 6/20
17/17 [=====] - 7s 408ms/step - loss: 0.6388 -
accuracy: 0.7932 - val_loss: 1.1764 - val_accuracy: 0.5909
Epoch 7/20
17/17 [=====] - 7s 400ms/step - loss: 0.4496 -
accuracy: 0.8482 - val_loss: 1.2553 - val_accuracy: 0.5076
Epoch 8/20
17/17 [=====] - 7s 397ms/step - loss: 0.3445 -
accuracy: 0.9013 - val_loss: 1.1815 - val_accuracy: 0.6364
Epoch 9/20
17/17 [=====] - 8s 439ms/step - loss: 0.2317 -
accuracy: 0.9469 - val_loss: 1.1707 - val_accuracy: 0.6364
Epoch 10/20
17/17 [=====] - 7s 404ms/step - loss: 0.1708 -
accuracy: 0.9696 - val_loss: 1.1616 - val_accuracy: 0.6364
Epoch 11/20
17/17 [=====] - 7s 403ms/step - loss: 0.1193 -
accuracy: 0.9772 - val_loss: 1.2087 - val_accuracy: 0.6061
Epoch 12/20
17/17 [=====] - 7s 402ms/step - loss: 0.0738 -
accuracy: 0.9924 - val_loss: 1.1887 - val_accuracy: 0.6591
Epoch 13/20
17/17 [=====] - 7s 438ms/step - loss: 0.0484 -
accuracy: 0.9981 - val_loss: 1.2497 - val_accuracy: 0.6364
Epoch 14/20
17/17 [=====] - 7s 438ms/step - loss: 0.0348 -
accuracy: 0.9981 - val_loss: 1.2645 - val_accuracy: 0.6439
Epoch 15/20
17/17 [=====] - 7s 404ms/step - loss: 0.0275 -
accuracy: 0.9981 - val_loss: 1.2805 - val_accuracy: 0.6591
Epoch 16/20
17/17 [=====] - 7s 397ms/step - loss: 0.0220 -
accuracy: 0.9981 - val_loss: 1.2952 - val_accuracy: 0.6515
Epoch 17/20
17/17 [=====] - 8s 439ms/step - loss: 0.0176 -
accuracy: 0.9981 - val_loss: 1.3430 - val_accuracy: 0.6439
Epoch 18/20
17/17 [=====] - 7s 424ms/step - loss: 0.0149 -
accuracy: 1.0000 - val_loss: 1.3086 - val_accuracy: 0.6515
Epoch 19/20

```
17/17 [=====] - 7s 398ms/step - loss: 0.0123 -  
accuracy: 1.0000 - val_loss: 1.3269 - val_accuracy: 0.6515  
Epoch 20/20  
17/17 [=====] - 7s 401ms/step - loss: 0.0106 -  
accuracy: 1.0000 - val_loss: 1.3416 - val_accuracy: 0.6515
```

```
[24]: accuracy = train_1.history['accuracy']  
val_accuracy = train_1.history['val_accuracy']  
  
epochs = range(len(accracy))  
  
fig = plt.figure(figsize=(12,8))  
plt.plot(epochs,accuracy,c="orange",label="Training")  
plt.plot(epochs,val_accuracy,c="yellow",label="Validation")  
plt.xlabel("Epochs")  
plt.ylabel("Accuracy")  
plt.legend()
```

```
[24]: <matplotlib.legend.Legend at 0x7c35d8365e10>
```



Rowan Banner ID - 916462025

Hence b) Train the CNN using 2 other number of filters: 4 and 16 for the convolution layer (i) with all other parameters unchanged

```
[25]: cnn_2 = Sequential([
    Conv2D(4, (3, 3), activation='relu', input_shape=(256, 256, 3)),
    MaxPool2D(2, 2),
    Flatten(),
    Dense(16, activation='relu'),
    Dense(4, activation='softmax')
])
cnn_2.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
cnn_2.summary()

train_2 = cnn_2.fit(train_data,
                    validation_data = test_data,
                    epochs = 20)

accuracy = train_2.history['accuracy']
val_accuracy = train_2.history['val_accuracy']

epochs = range(len(accuracy))

fig = plt.figure(figsize=(12,8))
plt.plot(epochs,accuracy,c="orange",label="Training")
plt.plot(epochs,val_accuracy,c="yellow",label="Validation")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 254, 254, 4)	112
max_pooling2d_1 (MaxPoolin g2D)	(None, 127, 127, 4)	0
flatten_1 (Flatten)	(None, 64516)	0
dense_2 (Dense)	(None, 16)	1032272
dense_3 (Dense)	(None, 4)	68

Total params: 1032452 (3.94 MB)
 Trainable params: 1032452 (3.94 MB)
 Non-trainable params: 0 (0.00 Byte)


```

-----
Epoch 1/20
17/17 [=====] - 8s 399ms/step - loss: 4.2745 -
accuracy: 0.3131 - val_loss: 2.5318 - val_accuracy: 0.2879
Epoch 2/20
17/17 [=====] - 6s 347ms/step - loss: 1.4827 -
accuracy: 0.3643 - val_loss: 1.3328 - val_accuracy: 0.3788
Epoch 3/20
17/17 [=====] - 6s 353ms/step - loss: 1.1432 -
accuracy: 0.4915 - val_loss: 1.2162 - val_accuracy: 0.3939
Epoch 4/20
17/17 [=====] - 6s 355ms/step - loss: 0.9790 -
accuracy: 0.6091 - val_loss: 1.0551 - val_accuracy: 0.5379
Epoch 5/20
17/17 [=====] - 6s 364ms/step - loss: 0.7697 -
accuracy: 0.7135 - val_loss: 1.1207 - val_accuracy: 0.5606
Epoch 6/20
17/17 [=====] - 7s 387ms/step - loss: 0.6573 -
accuracy: 0.7647 - val_loss: 1.1818 - val_accuracy: 0.5758
Epoch 7/20
17/17 [=====] - 6s 356ms/step - loss: 0.5208 -
accuracy: 0.8254 - val_loss: 1.0963 - val_accuracy: 0.5833
Epoch 8/20
17/17 [=====] - 6s 352ms/step - loss: 0.4198 -
accuracy: 0.8899 - val_loss: 1.1906 - val_accuracy: 0.5682
Epoch 9/20
17/17 [=====] - 6s 354ms/step - loss: 0.3960 -
accuracy: 0.8956 - val_loss: 1.3541 - val_accuracy: 0.5303
Epoch 10/20
17/17 [=====] - 6s 362ms/step - loss: 0.3809 -
accuracy: 0.8843 - val_loss: 1.2113 - val_accuracy: 0.6136
Epoch 11/20
17/17 [=====] - 6s 357ms/step - loss: 0.2910 -
accuracy: 0.9450 - val_loss: 1.1751 - val_accuracy: 0.6212
Epoch 12/20
17/17 [=====] - 6s 360ms/step - loss: 0.2267 -
accuracy: 0.9602 - val_loss: 1.1755 - val_accuracy: 0.6364
Epoch 13/20
17/17 [=====] - 6s 369ms/step - loss: 0.2292 -
accuracy: 0.9564 - val_loss: 1.2655 - val_accuracy: 0.5758
Epoch 14/20
17/17 [=====] - 7s 393ms/step - loss: 0.2311 -
accuracy: 0.9488 - val_loss: 1.2992 - val_accuracy: 0.5758
Epoch 15/20
17/17 [=====] - 6s 353ms/step - loss: 0.2067 -
accuracy: 0.9545 - val_loss: 1.2051 - val_accuracy: 0.6136
Epoch 16/20
17/17 [=====] - 6s 356ms/step - loss: 0.1605 -


```

```

accuracy: 0.9696 - val_loss: 1.2384 - val_accuracy: 0.6439
Epoch 17/20
17/17 [=====] - 6s 359ms/step - loss: 0.1400 -
accuracy: 0.9772 - val_loss: 1.2334 - val_accuracy: 0.6212
Epoch 18/20
17/17 [=====] - 6s 352ms/step - loss: 0.1347 -
accuracy: 0.9753 - val_loss: 1.2296 - val_accuracy: 0.6439
Epoch 19/20
17/17 [=====] - 6s 354ms/step - loss: 0.1212 -
accuracy: 0.9810 - val_loss: 1.2464 - val_accuracy: 0.6288
Epoch 20/20
17/17 [=====] - 7s 391ms/step - loss: 0.1064 -
accuracy: 0.9772 - val_loss: 1.2823 - val_accuracy: 0.6288

```

[25]: <matplotlib.legend.Legend at 0x7c35bcdaf2b0>



```

[26]: cnn_3 = Sequential([
    Conv2D(16, (3, 3), activation='relu', input_shape=(256, 256, 3)),
    MaxPool2D(2, 2),
    Flatten(),
    Dense(16, activation='relu'),
    Dense(4, activation='softmax')
])

```

```

cnn_3.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
cnn_3.summary()

train_3 = cnn_3.fit(train_data,
                   validation_data = test_data,
                   epochs = 20)

accuracy = train_3.history['accuracy']
val_accuracy = train_3.history['val_accuracy']

epochs = range(len(accuracy))

fig = plt.figure(figsize=(12,8))
plt.plot(epochs,accuracy,c="orange",label="Training")
plt.plot(epochs,val_accuracy,c="yellow",label="Validation")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 254, 254, 16)	448
max_pooling2d_2 (MaxPooling2D)	(None, 127, 127, 16)	0
flatten_2 (Flatten)	(None, 258064)	0
dense_4 (Dense)	(None, 16)	4129040
dense_5 (Dense)	(None, 4)	68

Total params: 4129556 (15.75 MB)

Trainable params: 4129556 (15.75 MB)

Non-trainable params: 0 (0.00 Byte)

Epoch 1/20

17/17 [=====] - 11s 590ms/step - loss: 9.6415 - accuracy: 0.2277 - val_loss: 2.3894 - val_accuracy: 0.3106


Epoch 2/20

17/17 [=====] - 10s 561ms/step - loss: 1.6883 - accuracy: 0.3454 - val_loss: 1.4788 - val_accuracy: 0.4318

Epoch 3/20
17/17 [=====] - 10s 609ms/step - loss: 1.3059 - accuracy: 0.3890 - val_loss: 1.1388 - val_accuracy: 0.5000
Epoch 4/20
17/17 [=====] - 10s 566ms/step - loss: 1.0517 - accuracy: 0.4782 - val_loss: 1.0898 - val_accuracy: 0.5076
Epoch 5/20
17/17 [=====] - 10s 563ms/step - loss: 0.9316 - accuracy: 0.5503 - val_loss: 1.1747 - val_accuracy: 0.5152
Epoch 6/20
17/17 [=====] - 10s 594ms/step - loss: 0.8632 - accuracy: 0.5958 - val_loss: 1.1144 - val_accuracy: 0.5000
Epoch 7/20
17/17 [=====] - 10s 569ms/step - loss: 0.8302 - accuracy: 0.5484 - val_loss: 1.2669 - val_accuracy: 0.4848
Epoch 8/20
17/17 [=====] - 10s 587ms/step - loss: 0.6383 - accuracy: 0.7306 - val_loss: 1.5191 - val_accuracy: 0.4924
Epoch 9/20
17/17 [=====] - 10s 569ms/step - loss: 0.5095 - accuracy: 0.8178 - val_loss: 1.1894 - val_accuracy: 0.5606
Epoch 10/20
17/17 [=====] - 10s 588ms/step - loss: 0.3937 - accuracy: 0.8861 - val_loss: 1.1720 - val_accuracy: 0.5833
Epoch 11/20
17/17 [=====] - 10s 565ms/step - loss: 0.3391 - accuracy: 0.9127 - val_loss: 1.3769 - val_accuracy: 0.5227
Epoch 12/20
17/17 [=====] - 10s 595ms/step - loss: 0.2662 - accuracy: 0.9222 - val_loss: 1.2057 - val_accuracy: 0.5758
Epoch 13/20
17/17 [=====] - 10s 576ms/step - loss: 0.2352 - accuracy: 0.9488 - val_loss: 1.4436 - val_accuracy: 0.5530
Epoch 14/20
17/17 [=====] - 10s 576ms/step - loss: 0.2130 - accuracy: 0.9336 - val_loss: 1.4518 - val_accuracy: 0.5606
Epoch 15/20
17/17 [=====] - 10s 603ms/step - loss: 0.1983 - accuracy: 0.9488 - val_loss: 1.2200 - val_accuracy: 0.6364
Epoch 16/20
17/17 [=====] - 10s 572ms/step - loss: 0.1809 - accuracy: 0.9545 - val_loss: 1.4571 - val_accuracy: 0.5833
Epoch 17/20
17/17 [=====] - 10s 614ms/step - loss: 0.1189 - accuracy: 0.9734 - val_loss: 1.6132 - val_accuracy: 0.5379
Epoch 18/20
17/17 [=====] - 10s 577ms/step - loss: 0.0852 - accuracy: 0.9848 - val_loss: 1.5012 - val_accuracy: 0.5682

```
Epoch 19/20
17/17 [=====] - 10s 604ms/step - loss: 0.0700 -
accuracy: 0.9962 - val_loss: 1.5129 - val_accuracy: 0.5682
Epoch 20/20
17/17 [=====] - 10s 582ms/step - loss: 0.0489 -
accuracy: 0.9962 - val_loss: 1.5818 - val_accuracy: 0.5682
```

[26]: <matplotlib.legend.Legend at 0x7c35bca1bd60>



On comparing the three models' accuracy graphs, `cnn_1` is mostly ideal since it has reached an stable state after certain epochs, yet training accuracy is very high compared to validation accuracy so it also might be overfitting, same goes for `cnn_2` model. whereas the `cnn_3` model is underfitting at first and overfitting later.